

STAT 443: Homework 3

Aronn Grant Laurel (21232475)

14 March, 2025

Question 1 ARMA Model

```
# this is where your R code goes

# Initialisation
data <- read.csv("employee_wages_total_industry.csv")


data.ts <- ts(data$employee_wages,
              start = c(1997, 1),
              frequency = 12)

# January 1997 to December 2018 as the training set for model fitting
data.ts.train <- window(data.ts,
                        start = c(1997, 1),
                        end = c(2018, 12),
                        frequency = 12)

# from January 2019 to December 2019 as the test set
data.ts.test <- window(data.ts,
                       start = c(2019, 1),
                       end = c(2019, 12),
                       frequency = 12)

plot(data.ts,
     main = "Canada Average Hourly Wage Rate",
     xlab = "Time (Year)",
     ylab = "Average Wage Rate ($)")
```

Canada Average Hourly Wage Rate



```
# LOESS Smoothing
# De-seasonalize Series
data.ts.ls <- stl(data.ts.train,
                  s.window = "periodic")

# From HW1
trend.ls <- data.ts.ls$time.series[, "trend"]
seas.ls <- data.ts.ls$time.series[, "seasonal"]
error.ls <- data.ts.ls$time.series[, "remainder"]

# ARMA Model fitting for Remainder Term
```


```
error.time.ls <- seq_along(error.ls)
error.ls.df <- data.frame("Remainder" = error.ls,
                          "Time" = error.time.ls)
```

From the time plot, we can observe clear upward trend meaning that our mean is not constant and therefore not stationary. We should apply an ARIMA series with $d > 0$ to transform the series into a stationary process, but in our question we will be fitting an ARMA model ($d = 0$).

```
# this is where your R code goes
error.ts <- ts(error.ls)


acf(error.ts,
     lag.max = 20,
     main = "Correlogram of Error Term")
```

Correlogram of Error Term



```
pacf(error.ts,  
      lag.max = 20,  
      main = "PACF of Error Term")
```

PACF of Error Term



Our ACF plot shows a ‘damped’ like sine wave and our PACF plot displays a negative cut off at lag 6 and 8 which we will use (p) for our ARMA model (a Pure AR Process) and assess the better fitted model.

```
# this is where your R code goes  
# Fitting our model  
model <- arima(error.ts, order = c(0, 0, 8))  
model
```

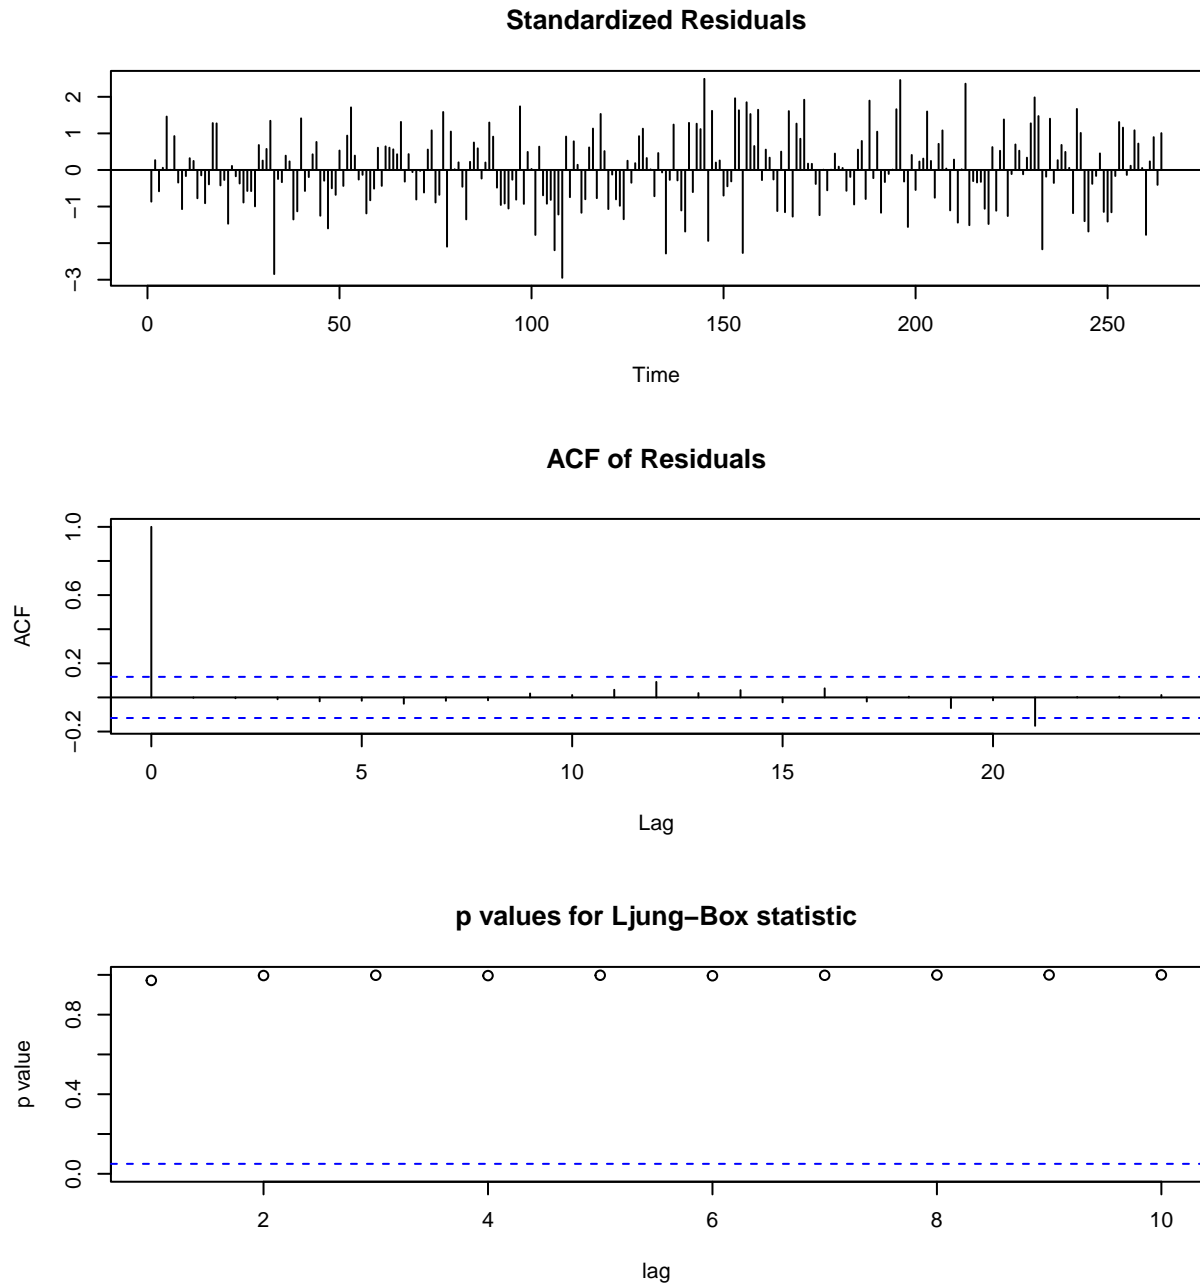
```
##  
## Call:  
## arima(x = error.ts, order = c(0, 0, 8))  
##
```

```
## Coefficients:
##      ma1      ma2      ma3      ma4      ma5      ma6      ma7      ma8
##      0.4853 0.1137 -0.1369 -0.2723 -0.3352 -0.3242 -0.1494 -0.1021
## s.e. 0.0615 0.0686 0.0650 0.0637 0.0642 0.0652 0.0698 0.0648
##      intercept
##      -6e-04
## s.e. 9e-04
##
## sigma^2 estimated as 0.002105: log likelihood = 438.26, aic = -856.52
```

```
model_2 <- arima(error.ts, order = c(0, 0, 6))
model_2
```

```
##
## Call:
## arima(x = error.ts, order = c(0, 0, 6))
##
## Coefficients:
##      ma1      ma2      ma3      ma4      ma5      ma6 intercept
##      0.4625 0.0485 -0.2004 -0.3228 -0.3510 -0.2731 -0.0007
## s.e. 0.0606 0.0607 0.0623 0.0675 0.0612 0.0572 0.0011
##
## sigma^2 estimated as 0.002144: log likelihood = 435.91, aic = -855.81
```

```
tsdiag(model)
```



After comparing the models of ARMA(8, 0) and ARMA(6, 0), I have decided on ARMA(8, 0) because it produces the smallest AIC value between the two models. Furthermore, its diagnostic plots show an arbitrary standardized residual plot with no obvious pattern. The ACF of residuals displays insignificant lag values aside from lag 0 which resembles a White Noise process along with non-significant p-values throughout our Ljung-Box Statistics plot. Hence, I believe our ARMA(8, 0) fitted model is the better fit for the remainder term.

The Fitted ARMA Model $X_t = -0.0006 + 0.4853Z_{t-1} + 0.1137Z_{t-2} - 0.1369Z_{t-3} - 0.2723Z_{t-4} - 0.3352Z_{t-5} - 0.3242Z_{t-6} - 0.1494Z_{t-7} - 0.1021Z_{t-8}$

```

# this is where your R code goes
# Remainder Component
remainder_forecast <- predict(model, n.ahead = 12)
remainder_lower <- remainder_forecast$pred - 1.96 * remainder_forecast$se
remainder_upper <- remainder_forecast$pred + 1.96 * remainder_forecast$se

# Trend Component
trend.time.ls <- seq_along(trend.ls)
trend.ls.df <- data.frame('trend'=trend.ls, 'time'=trend.time.ls)
trend_model.ls <- lm(trend ~ time, data=trend.ls.df)
trend_forecast_times <- max(trend.time.ls) + 1:12
trend_forecast <- predict(trend_model.ls, newdata=data.frame('time'=trend_forecast_times))

# Seasonal Component
seasonal_forecast <- seas.ls[1:12]

final_forecast <- trend_forecast + seasonal_forecast + remainder_forecast$pred

forecast_ts <- ts(final_forecast,
                  start=c(2019, 1),
                  frequency=12)

# Prediction Interval
lower_pi <- trend_forecast + seasonal_forecast + remainder_lower
upper_pi <- trend_forecast + seasonal_forecast + remainder_upper

lower_ts <- ts(lower_pi, start=c(2019, 1), frequency=12)
upper_ts <- ts(upper_pi, start=c(2019, 1), frequency=12)

# Prediction Table (From Lab 9)
forecast_table <- data.frame(Month = month.abb[1:12],
                             True_Value = data.ts.test,
                             Forecast = forecast_ts,
                             Lower_CI = lower_ts,
                             Upper_CI = upper_ts
                             )

kable(forecast_table,
      caption = "Monthly Forecasts of Average Hourly Wage Rates (2019)")

```

Table 1: Monthly Forecasts of Average Hourly Wage Rates (2019)

Month	True_Value	Forecast	Lower_CI	Upper_CI
Jan	28.14	28.14834	28.05843	28.23826
Feb	28.17	28.14567	28.04573	28.24562
Mar	28.14	28.16560	28.06513	28.26606
Apr	28.19	28.21311	28.11189	28.31432
May	28.12	28.08552	27.98139	28.18966
Jun	28.26	28.03917	27.93076	28.14758
Jul	28.13	27.93891	27.82665	28.05117
Aug	28.09	28.04496	27.93190	28.15802
Sep	28.62	28.45544	28.34200	28.56888

Month	True_Value	Forecast	Lower_CI	Upper_CI
Oct	28.65	28.49517	28.38173	28.60861
Nov	28.65	28.55944	28.44601	28.67288
Dec	28.63	28.57778	28.46434	28.69122

```

# this is where your R code goes
# From Lab 8 Prediction plot
plot(data.ts.test,
     main = "Average Hourly Wage Rates in 2019: Actual vs Predicted",
     ylab = "Hourly Wage Rates (CAD)",
     ylim = c(min(lower_pi)-0.2, max(upper_pi)+0.2),
     col = "black",
     type = "b",
     pch = 19)

# Prediction
lines(forecast_ts,
     col = "blue",
     type = "b",
     pch = 19)

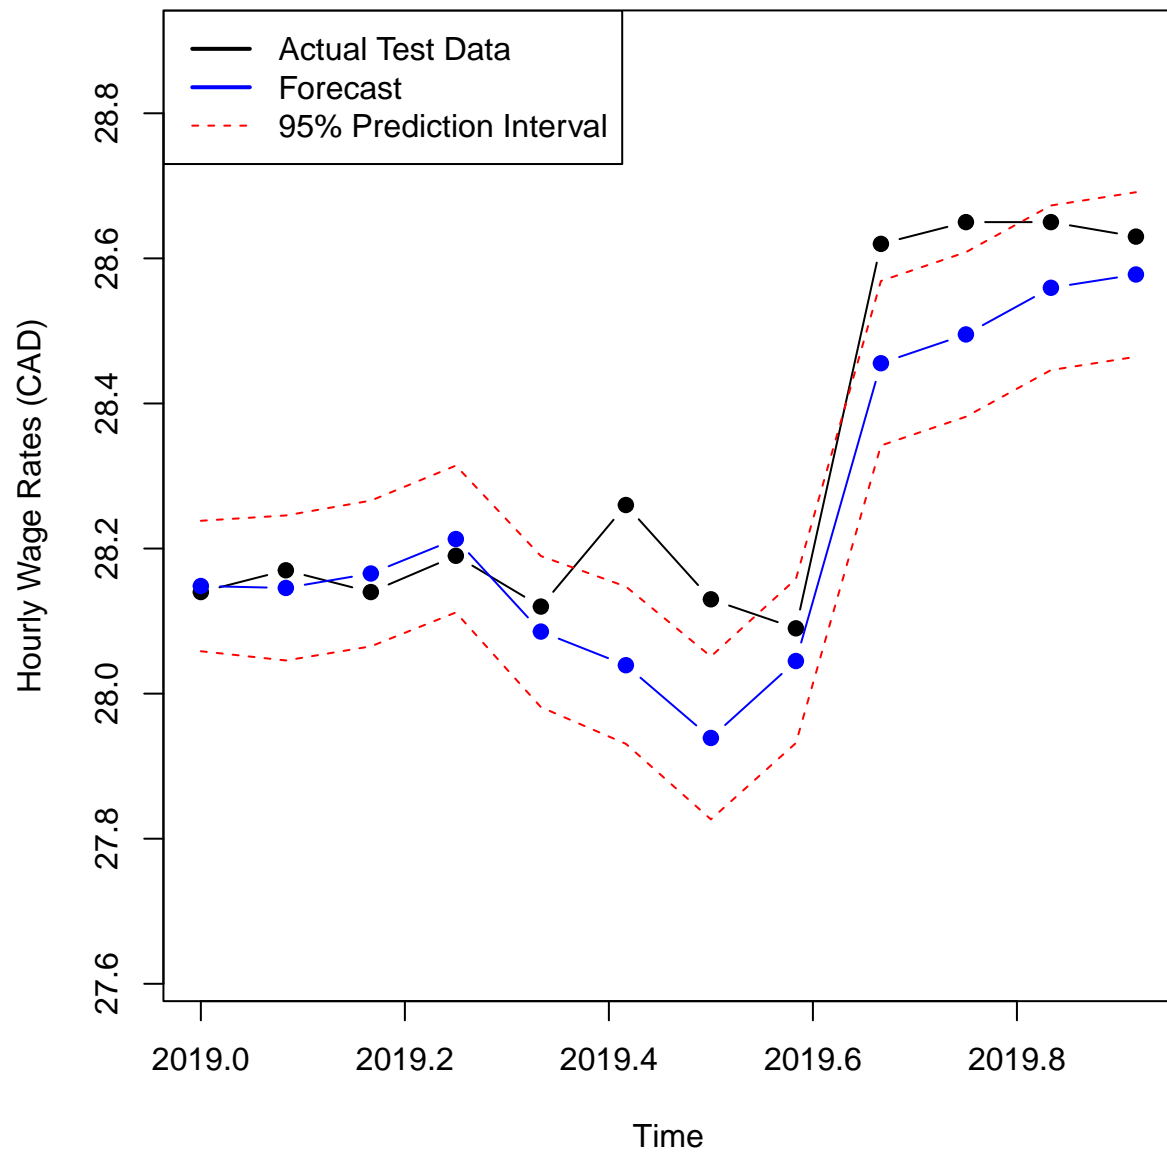
# Prediction interval
lines(lower_ts,
     type = "l",
     col = "red",
     lty = 2)

lines(upper_ts,
     type = "l",
     col = "red",
     lty = 2)

legend("topleft",
     legend=c("Actual Test Data", "Forecast", "95% Prediction Interval"),
     col=c("black", "blue", "red"),
     lty=c(1, 1, 2),
     lwd=c(2, 2, 1))

```


Average Hourly Wage Rates in 2019: Actual vs Predicted



Question 2 Box-Jenkins

```
# this is where your R code goes
plot(data.ts.train, main = "Time Series plot for Average Hourly Wage Rate")
```


Time Series plot for Average Hourly Wage Rate



Based on our time series plot, we can see a strong positive slope / trend, hence the mean is changing over time and we may consider it as non-stationary. Hence, I will be removing its trend and seasonal components to make the series stationary.

```
# this is where your R code goes  
acf(data.ts.train, main="ACF of Original Series")
```

ACF of Original Series




Now, we will be removing the linear trend by using the diff function with difference = 1 and remove its seasonality by differencing over 1 full season (12 months)

```
# FROM LAB 4
# Trend Differencing
diff_1 <- diff(data.ts.train, differences = 1)


plot(diff_1,
     main="First-Order Differenced Time Series",
     xlab = "Time", ylab = "Differenced Hourly Wage Rate")
```

First-Order Differenced Time Series




```
acf(ts(diff_1), main = "ACF of First Order Differenced Series")
```

ACF of First Order Differenced Series



```
pacf(ts(diff_1), main = "PACF of First Order Differenced Series")
```


PACF of First Order Differenced Series



```
# Seasonality Differencing
seasonal_diff <- diff(diff_1, lag = frequency(data.ts.train))


plot(seasonal_diff,
     main = "Seasonally Differenced Series",
     xlab = "Time", ylab = "Differenced Hourly Wage Rate"
)
```

Seasonally Differenced Series




```
acf(ts(seasonal_diff), lag.max = 30,  
    main="ACF of Seasonally Differenced Time Series")
```


ACF of Seasonally Differenced Time Series



```
pacf(ts(seasonal_diff), lag.max = 30,  
     main="PACF of Seasonally Differenced Time Series")
```

PACF of Seasonally Differenced Time Series

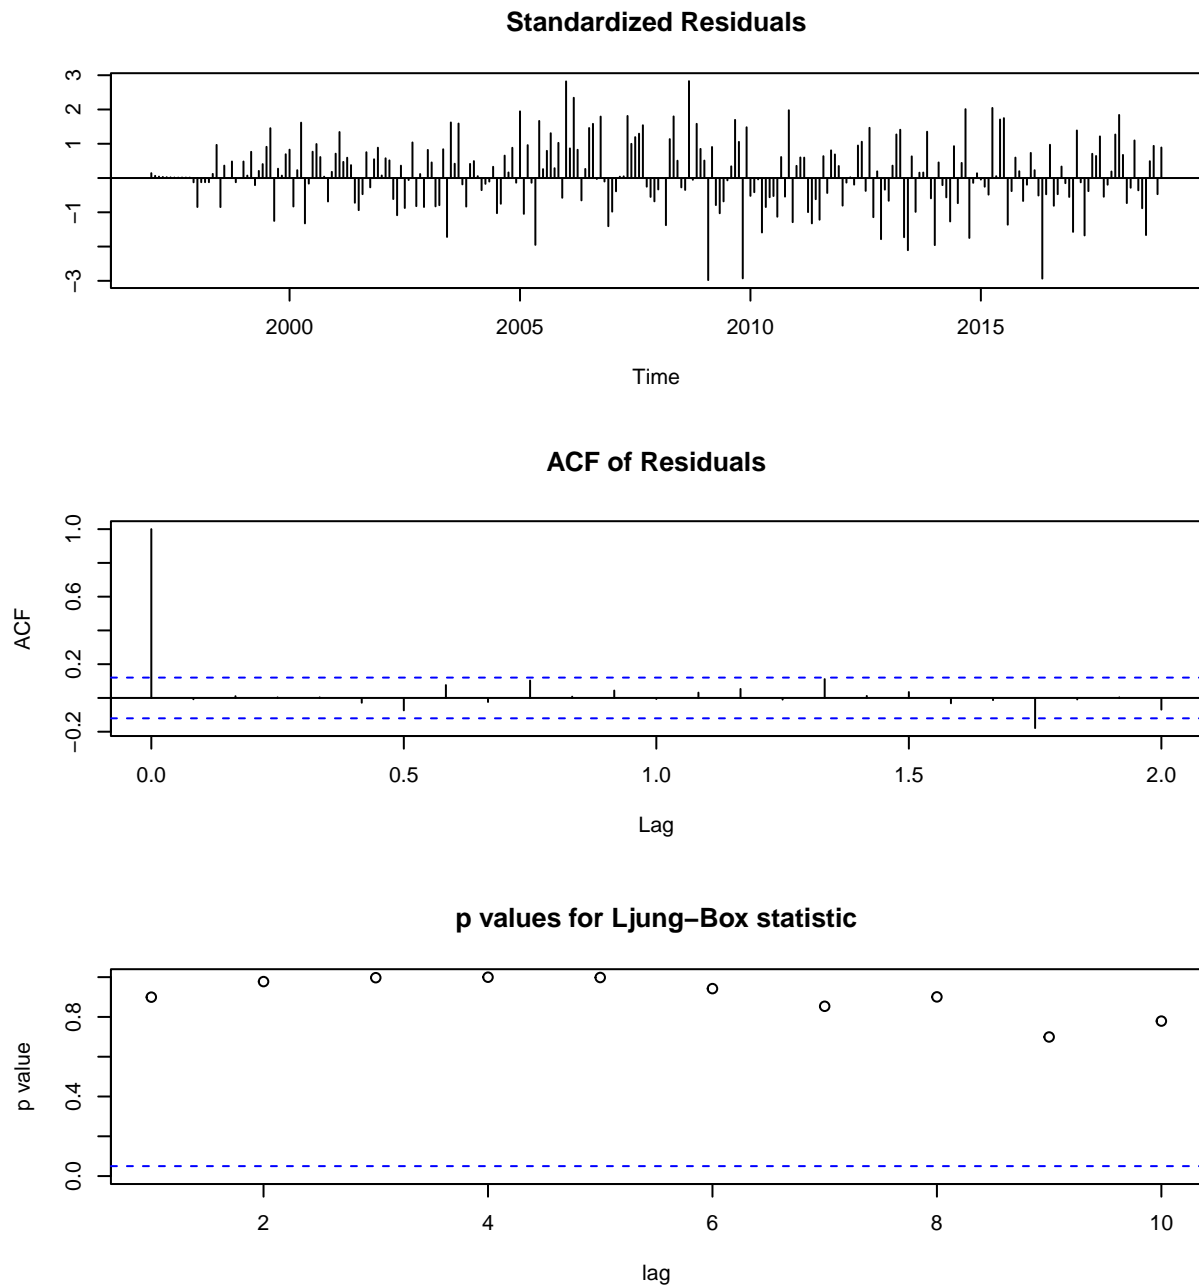


After differencing, it was quite difficult to determine the perfect cut off on our acf and pacf plots after our differencing and seasonal order. Hence, I will be using `auto.arima()` function to find the best ARIMA model for our time series training dataset.

```
# this is where your R code goes
arima_selection <- auto.arima(data.ts.train, seasonal = TRUE)
summary(arima_selection)
```

```
## Series: data.ts.train
## ARIMA(0,1,0)(0,1,1)[12]
##
## Coefficients:
##      sma1
##      -0.8247
## s.e.    0.0506
##
## sigma^2 = 0.004067: log likelihood = 328.74
## AIC=-653.48  AICc=-653.44  BIC=-646.43
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.005726408 0.06206067 0.04806889 0.03164616 0.2255326 0.08427272
##              ACF1
## Training set -0.007736081
```

```
tsdiag(arima_selection)
```



I believe that `SARIMA(0,1,0)*(0,1,1)_12` would fit the data best based on its small AIC value. Furthermore, the diagnostic plot suggests a good fit because the standardized residuals has no pattern, the acf of the residuals is a white noise process, and the p-values for Ljung-Box are insignificant.

```
# this is where your R code goes  
# Forecasting  
sarima_forecast <- predict(arima_selection, n.ahead=12)  
  
sarima_pred <- sarima_forecast$pred
```

```

sarima_forecast_lower <- sarima_pred - 1.96 * sarima_forecast$se
sarima_forecast_upper <- sarima_pred + 1.96 * sarima_forecast$se

sarima_forecast_ts <- ts(sarima_pred, start=c(2019, 1), frequency=12)
sarima_lower_ts <- ts(sarima_forecast_lower, start=c(2019, 1), frequency=12)
sarima_upper_ts <- ts(sarima_forecast_upper, start=c(2019, 1), frequency=12)

# Prediction Table (From Lab 9)
forecast_table <- data.frame(Month = month.abb[1:12],
                             True_Value = data.ts.test,
                             Forecast = sarima_forecast_ts,
                             Lower_CI = sarima_lower_ts,
                             Upper_CI = sarima_upper_ts
                             )

# Graph
plot(data.ts.test,
     main = "SARIMA Average Hourly Wage Rates in 2019: Actual vs Predicted",
     ylab = "Hourly Wage Rates (CAD)",
     ylim = c(min(sarima_lower_ts) - 0.2, max(sarima_upper_ts) + 0.2),
     col = "black",
     type = "b",
     pch = 19)

lines(sarima_forecast_ts,
     col="blue",
     type = "b",
     pch = 19)


lines(sarima_lower_ts,
     col="red",
     lty = 2)

lines(sarima_upper_ts,
     col="red",
     lty = 2)

legend("topleft",
     legend=c("Actual Test Data", "Forecast", "95% Prediction Interval"),
     col=c("black", "blue", "red"),
     lty=c(1, 1, 2),
     lwd=c(2, 2, 1))

```

SARIMA Average Hourly Wage Rates in 2019: Actual vs Predicted




Question 3 Holt-Winters

Looking back into our time series plot, we can observe a roughly constant seasonal variation throughout the series. Hence, I will be using an additive model for Holt-Winters' Exponential Smoothing Fit.

```
# From Lab 8
hw_model <- HoltWinters(data.ts.train, seasonal = "additive")

plot(hw_model, main="Holt-Winters' Exponential Smoothing Fit")
```

Holt-Winters' Exponential Smoothing Fit



```
hw_predictions <- predict(hw_model,
                          n.ahead = 12,
                          prediction.interval = TRUE)

hw_fit <- hw_predictions[, "fit"]
hw_lower <- hw_predictions[, "lwr"]
hw_upper <- hw_predictions[, "upr"]

hw_forecast_ts <- ts(hw_fit, start = c(2019, 1), frequency = 12)

hw_lower_ts <- ts(hw_lower, start = c(2019, 1), frequency = 12)
hw_upper_ts <- ts(hw_upper, start = c(2019, 1), frequency = 12)

# Prediction Table (From Lab 9)
hw_forecast_table <- data.frame(Month = month.abb[1:12],
                                True_Value = data.ts.test,
                                Forecast = hw_forecast_ts,
                                Lower_CI = hw_lower_ts,
                                Upper_CI = hw_upper_ts
                                )

kable(forecast_table,
      caption = "Monthly Forecasts of Average Hourly Wage Rates (2019)")
```

Table 2: Monthly Forecasts of Average Hourly Wage Rates (2019)

Month	True_Value	Forecast	Lower_CI	Upper_CI
Jan	28.14	28.07322	27.94821	28.19822
Feb	28.17	28.09777	27.92098	28.27455
Mar	28.14	28.14620	27.92969	28.36271
Apr	28.19	28.18110	27.93110	28.43111
May	28.12	28.01346	27.73394	28.29297
Jun	28.26	27.96230	27.65610	28.26849
Jul	28.13	27.84863	27.51790	28.17935
Aug	28.09	27.93615	27.58258	28.28971
Sep	28.62	28.36033	27.98532	28.73533
Oct	28.65	28.41103	28.01574	28.80632
Nov	28.65	28.47580	28.06121	28.89038
Dec	28.63	28.49898	28.06596	28.93201

```


plot(data.ts.test,
     main="Holt-Winters' Average Hourly Wage Rates: Prediction vs Actual (2019)",
     ylim=range(c(data.ts.test, hw_lower, hw_upper)),
     ylab="Wage Rates (CAD)",
     xlab="Time",
     col="black")

lines(hw_forecast_ts, col="blue", lwd=2)
lines(hw_lower_ts, col="red", lty=2)
lines(hw_upper_ts, col="red", lty=2)

legend("topleft",
     legend=c("Actual Test Data", "Holt-Winters Forecast", "95% Prediction Interval"),
     col=c("black", "blue", "red"),
     lty=c(1, 1, 2))

```

Holt-Winters' Average Hourly Wage Rates: Prediction vs Actual (201




Question 4 Model Comparison

```
# GRAPH
plot(data.ts.test,
     col = "black",
     pch = 19,
     type = "b",
     xlab="Month (2019)",
     ylab="Average Hourly Wage Rate (CAD)",
     main="Forecasting Methods vs True Value",
     ylim=range(c(data.ts.test, forecast_ts, sarima_forecast_ts, hw_forecast_ts))
)

lines(forecast_ts, col="red", lwd=2)
lines(sarima_forecast_ts, col="green", lwd=2)
lines(hw_forecast_ts, col="blue", lwd=2)

legend("topleft",
     legend=c("Test Data", "ARMA Model", "SARIMA Model", "Holt-Winters Forecast"),
     col=c("black", "red", "green", "blue"),
     lty=c(1, 1, 2))
```


Forecasting Methods vs True Value



```
# Table
arma_mspe <- mean((data.ts.test - forecast_ts)^2)
sarima_mspe <- mean((data.ts.test - sarima_forecast_ts)^2)
hw_mspe <- mean((data.ts.test - hw_forecast_ts)^2)

table_compare <- data.frame(Model = c("ARMA", "SARIMA", "Holt-Winters"),
                             MSPE = c(arma_mspe, sarima_mspe, hw_mspe))
kable(table_compare,
      caption = "Forecast Models Comparison w/ Mean Squared Prediction Error")
```

Table 3: Forecast Models Comparison w/ Mean Squared Prediction Error

Model	MSPE
ARMA	0.0126940
SARIMA	0.0320554
Holt-Winters	0.0327192

Based on our gathered forecasts, the true values for are mostly within our prediction intervals but we get ARMA (seasonal decomposition) as the best model amongst them all because it has the smallest MSPE.